

# An introduction to AI and Neural Networks

@ENSPIMA – Bordeaux INP

Jean-Luc.Charles@ENSAM.EU



September 2022



# Welcome to the course

## "An introduction to AI & Neural Networks"



An introduction to get familiarised with...

- Machine learning
- Training & operating Artificial Neural Networks
- Tensorflow & keras Python modules
- Applications to Predictive Maintenance.

# Welcome to the course

## "An introduction to AI & Neural Networks"



### An introduction to get familiarised with...

- Machine learning
- Training & operating Artificial Neural Networks
- Tensorflow & keras Python modules
- Applications to Predictive Maintenance.



### Ressources

- 3 hours of lecture and  $4 \times$  practical work Python sessions (4 x 3h) on **your laptop**
- Dedicated github repository with all the course material (PDF, notebooks, videos...)

# Practical Work: $4 \times 3h$

## Self training: Wake up your Python !

- Start with the 2 notebooks

`Wake_up_your_Python-part1.ipynb` and `...part2.ipynb`

# Practical Work: $4 \times 3h$

## Self training: Wake up your Python !

- Start with the 2 notebooks  
`Wake_up_your_Python-part1.ipynb` and `...part2.ipynb`

## Self training: AI & Machine Learning

- 3 *notebooks* `ML1_MNIST_en.ipynb`,  
`ML2_DNN_part1_en.ipynb` and `part2` target the skills:
  - load and pre-process MNIST images
  - build a **dense** neural network with `tensorflow` & `keras`
  - train the network to recognize MNIST images
  - evaluate and operate the trained network.

# Practical Work: $4 \times 3h$

## Self training: Wake up your Python !

- Start with the 2 notebooks  
`Wake_up_your_Python-part1.ipynb` and `...part2.ipynb`

## Self training: AI & Machine Learning

- 3 *notebooks* `ML1_MNIST_en.ipynb`,  
`ML2_DNN_part1_en.ipynb` and `part2` target the skills:
  - load and pre-process MNIST images
  - build a **dense** neural network with `tensorflow` & `keras`
  - train the network to recognize MNIST images
  - evaluate and operate the trained network.

## Mini-project: application to Predictive Maintenance

- use your skills to process a case of predictive maintenance with a dense neural network...

# The historical way...

## ARTIFICIAL INTELLIGENCE

Early artificial intelligence stirs excitement.



## MACHINE LEARNING

Machine learning begins to flourish.



## DEEP LEARNING

Deep learning breakthroughs drive AI boom.



1950's

1960's

1970's

1980's

1990's

2000's

2010's

(from : [developer.nvidia.com/deep-learning](https://developer.nvidia.com/deep-learning))

# Artificial Intelligence ?

**Artificial Intelligence**<sup>1</sup>: remains an ambiguous term with multiple definitions varying with time:

- *"...the science of making computers do things that require intelligence when done by humans."* [Alan Turing, 1940](#)
- *"the field of study that gives computers the ability to learn without being explicitly programmed."* [Arthur Samuel, 1960](#)
- *"A computer program is said to learn from experience  $E$  with respect to some class of tasks  $T$  and performance measure  $P$ , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ ."* [Tom Mitchell, 1997](#)
- Notion of *intelligent agent* or *rational agent*  
*"...agent that acts in such a way as to reach the best solution or, in an uncertain environment, the best predictable solution."* [Stuart Russel, Peter Norvig, "Intelligence Artificielle" 2015](#)

---

<sup>1</sup> first used in 1956 by [John McCarthy](#), researcher at Stanford during the Dartmouth conference



# Artificial Intelligences ?

**Strong AI**

**Weak AI**

**General AI**

**Narrow AI**

# Artificial Intelligences ?

## Strong AI

- Build systems that think exactly the same way that people do.
- Try also to explain how humans think...
- We are not yet here...

## Weak AI

## General AI

## Narrow AI

# Artificial Intelligences ?

## Strong AI

- Build systems that think exactly the same way that people do.
- Try also to explain how humans think...
- We are not yet here...

## Weak AI

- Build systems that can behave like humans.
- The results will tell us nothing about how humans think.
- We already are there... We use it every day!  
(anti-spam, facial or voice recognition, language translation...)

## General AI

## Narrow AI

# Artificial Intelligences ?

## Strong AI

- Build systems that think exactly the same way that people do.
- Try also to explain how humans think...
- We are not yet here...

## Weak AI

- Build systems that can behave like humans.
- The results will tell us nothing about how humans think.
- We already are there... We use it every day!  
(anti-spam, facial or voice recognition, language translation...)

## General AI

- AI systems designed for the ability to reason in general.

## Narrow AI

- AI systems designed for specific tasks.

# Artificial Intelligence

- Runs in much of our present technology (smartphone apps...)
- Powered by rapid advances in data storage, computer processing power
- Powered by free dataset acces via Internet and code publishing as open source environments
- Rate of acceleration is already astounding
- Will likeky shape our future more powerfully than any other innovation this century.

# Some AI famous dates

- May 11, 1997: the IBM computer [Deep Blue](#) defeated Gary Kasparov at chess.  
Today, Kasparov says: *“computer that defeated me at chess is no more intelligent than an alarm clock”*  
The ability to defeat a human is no more a criteria for defining AI.
- 2011: IBM’s Watson computer wins television game show Jeopardy, defeating legendary human champions.
- 2015: Google [Deepmind](#) developped an agent that surpassed human performances at 49 Atari games
- 2016: Google DeepMind’s [AlphaGo](#) defeats Go champion Lee Sedol.

# Machine Learning and AI

Page from [medium.com/machine-learning-for-humans/...](https://medium.com/machine-learning-for-humans/)

## Machine learning $\subseteq$ artificial intelligence

### ARTIFICIAL INTELLIGENCE

Design an intelligent agent that perceives its environment and makes decisions to maximize chances of achieving its goal.

Subfields: vision, robotics, machine learning, natural language processing, planning, ...

### MACHINE LEARNING

Gives "computers the ability to learn without being explicitly programmed" (Arthur Samuel, 1959)

#### SUPERVISED LEARNING

Classification, regression

#### UNSUPERVISED LEARNING

Clustering, dimensionality  
reduction, recommendation

#### REINFORCEMENT LEARNING

Reward maximization

# Branches of Machine Learning

## Supervised learning

Needs data and labels...

- **Classification**

- Images classification
- Objects detection
- speech recognition...

- **Regression**

- predict a value...

- **Anomaly detection**

- Spam detection
- Manufacturing: finding known (learned) defects
- Weather prediction
- Diseases classification...



# Branches of Machine Learning

## Unsupervised learning

Needs only data...

- **Clustering** – non labelled data **Grouping**
  - Data mining, web data grouping, news grouping...
  - Market segmentation
  - DNA grouping
  - Astronomical data analysis...
- **Anomaly Detection**
  - Fraud detection
  - Manufacturing: finding defects even new ones
  - Monitoring activity: detect abnormal activity (failure, hacker, fraud...)
  - Fake account on Internet...
- **Dimensionality reduction**
  - Compress data using fewer numbers...

# Branches of Machine Learning

## Reinforcement learning

An agent learns to drive an environment...

- **Reward maximisation**
  - ...
- **Control/command**
  - Controlling robots, drones...
  - Factory optimization
  - Financial (stock) trading...
- **Decision making**
  - games (video games)
  - financial analysis...

# Machine Learning approaches

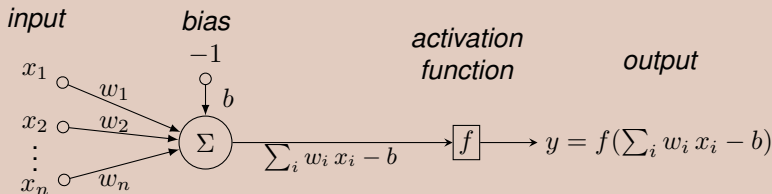
Several approaches/technics can be used to design *Machine Learning* algorithms:

- Genetic programming
- Bayesian inference
- Fuzzy logic
- Neural Networks
- ...

The following deals only with **Artificial Neural Networks**.

# The artificial neuron

## The computer model of the artificial neuron

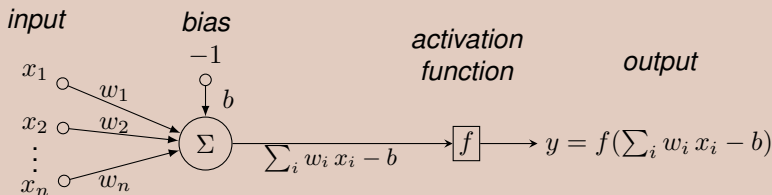


An **artificial neuron**:

- receives the input data  $(x_i)_{i=1..n}$  affected by the **weights**  $(w_i)_{i=1..n}$  (*weights*)

# The artificial neuron

## The computer model of the artificial neuron

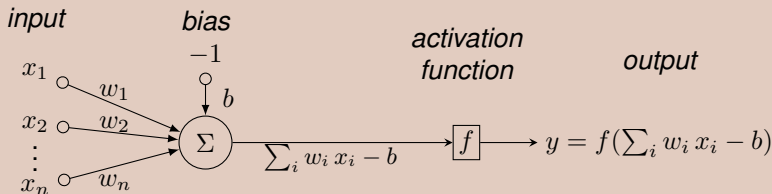


An **artificial neuron**:

- receives the input data  $(x_i)_{i=1..n}$  affected by the **weights**  $(w_i)_{i=1..n}$  (*weights*)
- calculates the **weighted sum** of its entries minus the bias  $\sum_i w_i x_i - b$

# The artificial neuron

## The computer model of the artificial neuron



An **artificial neuron**:

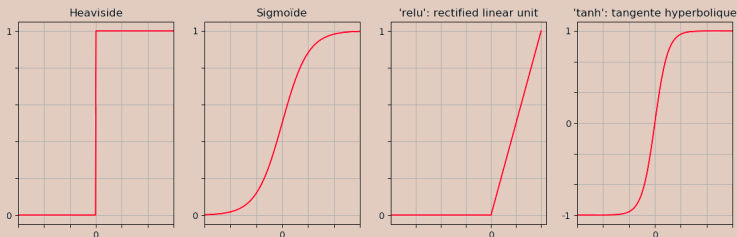
- receives the input data  $(x_i)_{i=1..n}$  affected by the **weights**  $(w_i)_{i=1..n}$  (*weights*)
- calculates the **weighted sum** of its entries minus the bias  $\sum_i w_i x_i - b$
- outputs a **activation**  $f(\sum_i w_i x_i - b)$ , computed with an activation function  $f$  (generally non-linear).

# Artificial neuron

The activation function of a neuron:

- introduces a non-linear behavior,
- sets the range of the neuron output, for example  $[-1, 1]$ ,  $[0, 1]$  or even  $[0, \infty[$ .

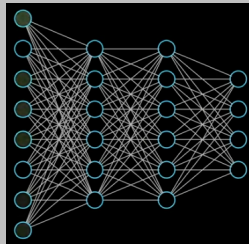
## Examples of often used activation functions



- The bias  $b$  sets the activation threshold of the neuron.

# Neural networks studied

- Neural networks are more or less complex assemblies of artificial neurons grouped by layers.



- Two architectures are very common:
  - The **Dense Neural Network** (DNN), simple, generalist, can perform greatly when well tuned.
  - The more complex **Convolutional Neural Network** (CNN), mainly specialized in image processing.



# A must example: training a Dense Network to classify the MNIST handwritten digit images

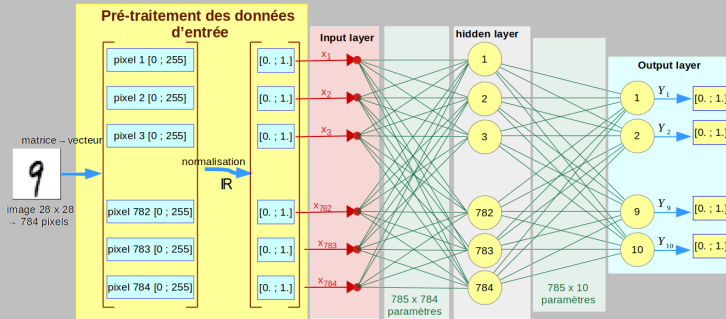
- **MNIST**: bank of 70000 **labeled images**  
(60000 training images and 10000 test images)



- grayscale images  $28 \times 28$  pixels.
- Scores with a dense networks can reach 98% success...
- State of the art for image recognition : **Convolutional Neural Networks** (CNN)  
[will not be covered in this course limited to dense networks]

# Dense Neural Network architecture

Each matrix  $28 \times 28 \rightsquigarrow$  normalized vector of 784 components  $\text{float} \in [0; 1]$ .



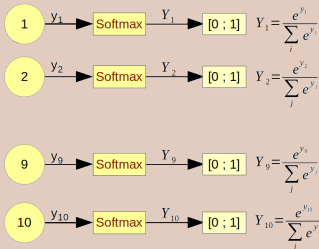
Structure of the network:

- An *Input layer* sets the size of network inputs to 784 values. It has no neurons.
- A *Hidden layer* of 784 neurons (we could have more, or less...), receives the input data. It is connected to the next layer.
- An *Output layer* of 10 neurons (1 neuron for each digit to be recognized).

# Activation functions

- In the intermediate layers the activation function *relu* often favors the learning of the network <sup>2</sup> algorithm.
- Classification (last layer) uses the *softmax* function:

## Activation function *softmax*



- The activation of neuron  $k$  is  $Y_k = e^{y_k} / \sum_i e^{y_i}$  with  $y_k = \sum_i \omega_i x_i - b$  calculated by the neuron  $k$ .
- The outputs of the neurons are interpreted as probabilities in the interval  $[0, 1]$ .

The neuron with the greatest probability (activation) gives the response of the network by its associated label.

<sup>2</sup> avoids the *vanishing gradient* that appears in the *back propagation*

## One-hot encoding of labels

Purpose: to put the image labels in the format of the network output

- Image labels: **integers** from 0 to 9.
- Network output: **vector of 10 float** in the interval  $[0,1]$  calculated by the *softmax* functions of the 10 output neurons.
- *one-hot* coding of an ordered collection of  $N$  unique elements:

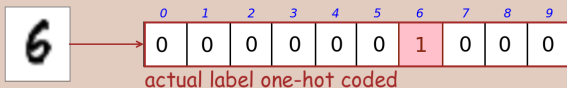
- each element is coded by a vector of  $N$  null components except one,
- the  $i$ th element  $\leadsto$  vector with a 1 for  $i$ th component.

chiffre	$Y_i'$ : vecteur <i>one-hot</i>
0	[1 0 0 0 0 0 0 0 0 0]
1	[0 1 0 0 0 0 0 0 0 0]
2	[0 0 1 0 0 0 0 0 0 0]
3	[0 0 0 1 0 0 0 0 0 0]
4	[0 0 0 0 1 0 0 0 0 0]
5	[0 0 0 0 0 1 0 0 0 0]
6	[0 0 0 0 0 0 1 0 0 0]
7	[0 0 0 0 0 0 0 1 0 0]
8	[0 0 0 0 0 0 0 0 1 0]
9	[0 0 0 0 0 0 0 0 0 1]

The *one-hot* encoding of labels '0' to '9' results in a 10-component vector, like the one computed by the neural network.

## Error function: *Cross entropy error*

- An image processed by the network  $\leadsto$  vector  $\hat{Y}$  of 10 float to compare to the *hot-one* encoding  $Y$  of the label of the image.
- We use the error (or loss) function *cross entropy* adapted to the coding *one-hot*:  $e(Y, \hat{Y}) = -\sum_i Y_i \log(\hat{Y}_i)$



cross entropy error:  $-\sum_i Y_i \log(\hat{Y}_i)$

computed probabilities

0.1	0.2	0.1	0.3	0.2	0.1	0.9	0.3	0.1	0.2
0	1	2	3	4	5	6	7	8	9

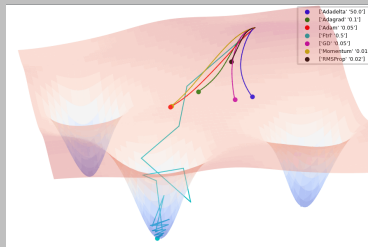
## Optimization and *Back Propagation*

- *Feed forward stage* : an optimization algorithm calculates the gradient of the loss function relative to network weights.
- *Back Propagation* : the BP algorithm **modifies** the weights of the network thanks to the gradient of the loss function, iterating from the last layer to the first layer.
- Examples of optimization algorithm used:
  - *Gradient Descent* (GD)
  - *Stochastic Gradient Descent* (SGD)
  - *Adam* (enhanced version of gradient descent)...

The module [tf.keras.optimizers](#) offers Python implementation of several optimization algorithms.

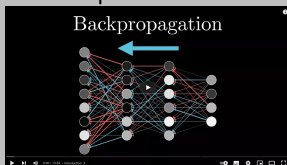
# Dense Neural Network

Visualization of gradient descent algorithm iterations for an ultra-simple loss function with only 2 variables:



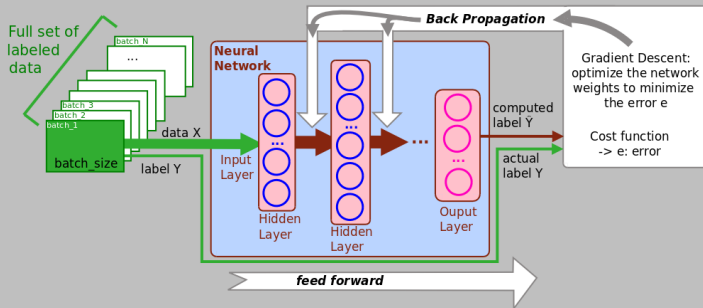
(source: [github.com/Jaewan-Yun/optimizer-visualization](https://github.com/Jaewan-Yun/optimizer-visualization))

*back propagation* algorithm explanation video:



# Supervised learning strategy

## Supervised learning : Feed Forward and Back Propagation

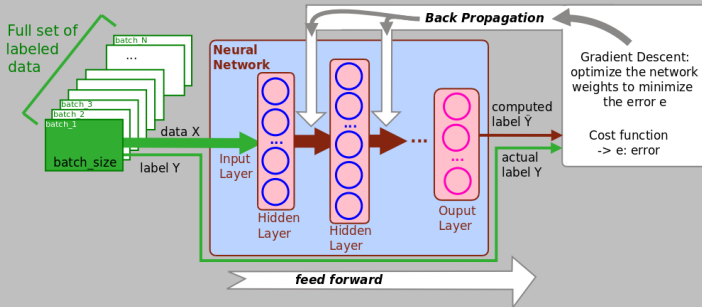


- The full data set is splitted in (mini) **batches** of size `batch_size`



# Supervised learning strategy

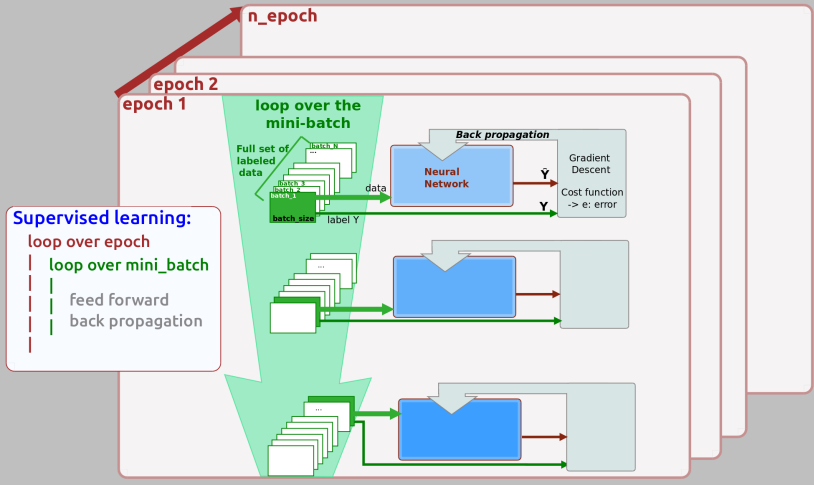
## Supervised learning : Feed Forward and Back Propagation



- The full data set is splitted in (mini) **batches** of size **batch\_size**
- After each batch has been fed forward: the *Back Propagation* algorithm modifies the weights of the network layer by layer to minimize the error  $e$ .

# Supervised learning strategy

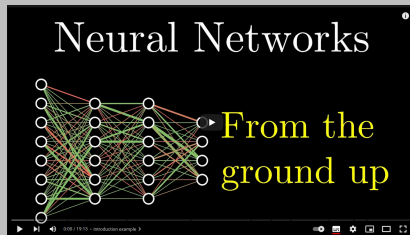
- The training over the full data set is repeated  $n\_epoch$  times....



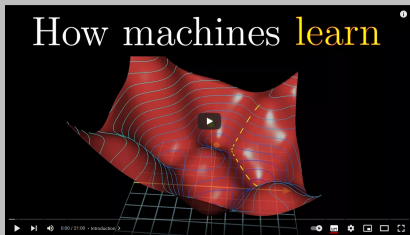
# Videos



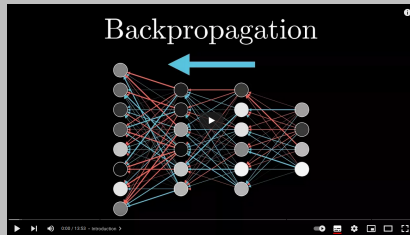
1/ Local: "Le deep learning - YouTube.webm"



2/ local : "But what is a neural network.webm"



3/ Local: "Gradient descent how neural networks learn.webm"



4/ Local: "What is backpropagation really doing .webm"

# References

- [1] *Artificial Intelligence: A Modern Approach (4th Edition)*, By Stuart Russell & Peter Norvig. Pearson, 2020. ISBN 978-0134610993. [aima.cs.berkeley.edu](http://aima.cs.berkeley.edu)

*Intelligence artificielle – Une approche moderne – 4e éd.*, By Stuart Russell & Peter Norvig. Translated by L. Miclet, F. Popineau, & C. Cadet. Paris: Pearson Education France, 2021. ISBN 978-2326002210.

- [2] *What is artificial intelligence (AI), and what is the difference between general AI and narrow AI?*, Kris Hammond, 2015  
[www.computerworld.com/article/2906336/what-is-artificial-intelligence.html](http://www.computerworld.com/article/2906336/what-is-artificial-intelligence.html)

- [3] *Stanford Encyclopedia of Philosophy*, [plato.stanford.edu/entries/artificial-intelligence](http://plato.stanford.edu/entries/artificial-intelligence)

- [4] *Deep Learning.*, Goodfellow, Ian; Bengio, Yoshua; Courville, Aaron (2016), MIT Pres, ISBN 9780262035613